

Name: Steven Phung

ID: 028023383

Professor: Ju Cheol Moon

Due: 4/7/22

### Assignment #8 Report

The open source code of YOLO V3 that I referenced and used for this assignment is called YOLOFace. I used the <https://github.com/nbishdev/yoloface> fork because it solved an index error that the main branch had. The pre-trained weights I used were the ones linked in the repo: yolov3-wider\_16000.weights at <http://shuoyang1213.me/WIDERFACE/>. I modified some of YOLOFace's files to have smoother output. For yoloface.py I commented out cv2 windows that showed the images since I just wanted it to run in the background and it not have to interrupt the system every few seconds. I also modified utils.py to write the coordinates of the bounding box to a file. Then I used these coordinates for image2vect.py to crop each image to only contain the face. Finally, it returned an embedding vector for the cropped image.

The program imageFinder.py takes a single command line input as the initial file to calculate distances for. The input is simply *file\_name* such as "000109.jpg" or "000308.jpg." The embedding vector is received from image2vect.py and then done similarly for the other 1,199 images. Then I implemented euclidean\_distance(vectorA, vectorB) function manually, it takes 2 vectors, for each vector's index it subtracts, then squares this summation, and adds it to a total. At the very end I square root the entire total and that is my calculated Euclidean distance. After this I had a list of hyper-parameter Tau's from 0-2 incrementing by 0.1. Based on the Tau values it would output which image files were recognized as the same celebrity of the input image. This was repeated for 9 other randomly selected unique celebrities. I am assuming 10 total unique so I

simply generated 9 others from a random sample which excludes the initial input celebrity.

Finally, precision and recall were calculated for each celebrity and each tau in the interval.

I am unsure about my results and I think it may have to do with my Euclidean distance calculations. I think results were off because some images were much larger or some images were much smaller than others, resulting in large differences even if they were similar celebrities. If I had time to go back and repeat this project I would consider implementing some function that would resize the cropped images to all be a specific dimension, which would in turn result in equal embedding vectors. Lastly, I expected this program to take more time to finish, but even in the end it took more time than I expected.

Output for precision and recall curves:



